

Predykcja ataku **phishingowego** w wiadomości e-mail za pomocą **nadzorowanego nauczania maszynowego**

Dataset:

- **Phishing Email Curated Datasets**
 - <https://zenodo.org/records/8339691>

Najważniejsze użyte **moduły**

- **pandas** - praca z *Data Framami*
- **numpy** - obliczenia
- **matplotlib.pyplot** - wizualizacja
- **sklearn** - wszelakie narzędzia do *Machine Learningu*

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import KFold, train_test_split, cross_val_score, GridS
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import StandardScaler
import sklearn.metrics as skm

import import_ipynb
import re
from unidecode import unidecode
from termcolor import colored

import warnings
warnings.simplefilter(action='ignore', category=UserWarning)
```

Wczytanie przygotowanego **data framu**

```
In [2]: learning_set = pd.read_csv('ML_DataFrame.csv', index_col=0)
print(learning_set.head(3))
```

	label	suspicious_words_subject	suspicious_words_body	sender_nums_count	\
0	1.0	0.0	2.0	0.0	
1	1.0	0.0	0.0	4.0	
2	1.0	0.0	4.0	0.0	

	sender_domain_num_count	sender_domain_length	urls_count	protocol	\
0	0.0	6.0	1.0	0.0	
1	0.0	6.0	1.0	0.0	
2	0.0	16.0	3.0	0.0	

	contains_ip	url_length	TLD_alpha	subdomain_level	slash_count	\
0	0.0	21.0	1.0	0.0	3.0	
1	0.0	25.0	1.0	1.0	2.0	
2	0.0	72.0	1.0	1.0	6.0	

	dots_count	hyphens_count	has_non_latin
0	1.0	0.0	0.0
1	2.0	0.0	0.0
2	4.0	0.0	0.0

Usunięcie wierszy z pustymi wartościami

```
In [3]: print(learning_set.isna().sum())
learning_set.dropna(inplace=True)
```

label	0
suspicious_words_subject	0
suspicious_words_body	0
sender_nums_count	0
sender_domain_num_count	0
sender_domain_length	0
urls_count	0
protocol	0
contains_ip	0
url_length	0
TLD_alpha	0
subdomain_level	84
slash_count	0
dots_count	0
hyphens_count	0
has_non_latin	0
dtype: int64	

Wybranie X i y

X:

- suspicious_words_subject
- suspicious_words_body
- sender_nums_count
- sender_domain_num_count
- sender_domain_length
- urls_count

- protocol
- contains_ip
- url_length
- TLD_alpha
- subdomain_level
- slash_count
- dots_count
- hyphens_count
- has_non_latin

y:

- label

```
In [4]: X = learning_set.loc[:, 'suspicious_words_subject':'has_non_latin'].values
y = learning_set.loc[:, 'label'].values
print(X.shape, y.shape)
```

(41574, 15) (41574,)

Rozdzielanie **X** i **y** na treningowe i testowe zestawy

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Skalujemy wartości **X-ów**

```
In [6]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [7]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(33259, 15) (8315, 15) (33259,) (8315,)

Szukanie najlepszych parametrów dla **KNeighborsClassifier**

```
In [8]: kf = KFold(n_splits=6, shuffle=True, random_state=42)
params = {
    'n_neighbors': np.arange(1, 15, 1),
    'weights': ['uniform', 'distance'],
    'p': [1, 2]
}
knn = KNeighborsClassifier()
knn_cv = GridSearchCV(knn, param_grid=params, cv=kf)
knn_cv.fit(X_train, y_train)

print('Najlepsze parametry dla KNeighborsClassifier:')
for p, val in knn_cv.best_params_.items():
    print('{}: {}'.format(p, val), end='\n')
print('Uzyskana precyzja: ', knn_cv.best_score_)
```

```
best_knn = knn_cv.best_estimator_  
test_accuracy = best_knn.score(X_test, y_test)  
print("Precyzja zestawu testowego:", test_accuracy)
```

Najlepsze parametry dla KNeighborsClassifier:

n_neighbors: 8

p: 1

weights: distance

Uzyskana precyzja: 0.8931714625269559

Precyzja zestawu testowego: 0.8939266386049308

Szukanie najlepszych parametrów dla *LogisticRegression*

```
In [ ]: from sklearn.linear_model import LogisticRegression  
kf = KFold(n_splits=6, shuffle=True, random_state=42)  
params = {  
    'penalty': [None, 'l2'],  
    'C': [0.001, 0.01, 0.1, 1, 10]  
}  
logreg_forest = LogisticRegression()  
logreg_cv = GridSearchCV(logreg_forest, param_grid=params, cv=kf)  
logreg_cv.fit(X_train, y_train)  
  
print('Najlepsze parametry dla LogisticRegression:')  
for p, val in logreg_cv.best_params_.items():  
    print('{}: {}'.format(p, val), end='\n')  
print('Uzyskana precyzja: ', logreg_cv.best_score_  
  
best_logreg = logreg_cv.best_estimator_  
test_accuracy = best_logreg.score(X_test, y_test)  
print("Precyzja zestawu testowego:", test_accuracy)
```

Szukanie najlepszych parametrów dla *DecisionTreeClassifier*

```
In [10]: kf = KFold(n_splits=6, shuffle=True, random_state=42)  
params = {  
    'criterion': ['gini', 'entropy'],  
    'max_depth': [None, 5, 10, 15],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4]  
}  
tree = DecisionTreeClassifier()  
tree_cv = GridSearchCV(tree, param_grid=params, cv=kf)  
tree_cv.fit(X_train, y_train)  
  
print('Najlepsze parametry dla DecisionTreeClassifier:')  
for p, val in tree_cv.best_params_.items():  
    print('{}: {}'.format(p, val), end='\n')  
print('Uzyskana precyzja: ', tree_cv.best_score_  
  
best_tree = tree_cv.best_estimator_  
test_accuracy = best_tree.score(X_test, y_test)  
print("Precyzja zestawu testowego:", test_accuracy)
```

Najlepsze parametry dla DecisionTreeClassifier:
criterion: gini
max_depth: None
min_samples_leaf: 1
min_samples_split: 10
Uzyskana precyzja: 0.8818364460390439
Precyzja zestawu testowego: 0.8817799158147925

Szukanie najlepszych parametrów dla *RandomForestClassifier*

```
In [11]: from sklearn.ensemble import RandomForestClassifier

kf = KFold(n_splits=6, shuffle=True, random_state=42)
params = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
rand_forest = RandomForestClassifier()
rand_forest_cv = GridSearchCV(rand_forest, param_grid=params, cv=kf)
rand_forest_cv.fit(X_train, y_train)

print('Najlepsze parametry dla RandomForestClassifier:')
for p, val in rand_forest_cv.best_params_.items():
    print('{}: {}'.format(p, val), end='\n')
print('Uzyskana precyzja: ', rand_forest_cv.best_score_)

best_rand_forest = rand_forest_cv.best_estimator_
test_accuracy = best_rand_forest.score(X_test, y_test)
print("Precyzja zestawu testowego:", test_accuracy)
```

Najlepsze parametry dla RandomForestClassifier:
max_depth: None
min_samples_leaf: 1
min_samples_split: 5
n_estimators: 100
Uzyskana precyzja: 0.9137373190683672
Precyzja zestawu testowego: 0.9128081779915814

Szukanie najlepszych parametrów dla *SVM Classifier*

```
In [12]: from sklearn.svm import SVC

kf = KFold(n_splits=6, shuffle=True, random_state=42)
params = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf'],
    'gamma': ['scale', 'auto', 0.1, 1]
}
svm = SVC()
svm_cv = GridSearchCV(svm, param_grid=params, cv=kf)
svm_cv.fit(X_train, y_train)

print('Najlepsze parametry dla SVM Classifier:')
```

```

for p, val in svm_cv.best_params_.items():
    print('{}: {}'.format(p, val), end='\n')
print('Uzyskana precyzja: ', svm_cv.best_score_)

best_svm = svm_cv.best_estimator_
test_accuracy = best_svm.score(X_test, y_test)
print("Precyzja zestawu testowego:", test_accuracy)

```

Najlepsze parametry dla SVM Classifier:

C: 1

gamma: 1

kernel: rbf

Uzyskana precyzja: 0.8840010675208222

Precyzja zestawu testowego: 0.8849067949488876

Ewaluacja modeli

```

In [13]: models = {
    'KNeighborsClassifier': KNeighborsClassifier(n_neighbors=8, p=1, weights='distance'),
    'LogisticRegression': LogisticRegression(C=1, penalty='l2'),
    'DecisionTreeClassifier': DecisionTreeClassifier(criterion='gini', max_depth=None),
    'RandomForestClassifier': RandomForestClassifier(max_depth=None, min_samples_leaf=10),
    'SVM Classifier': SVC(C=10, gamma=1, kernel='rbf')
}

results = []
for model in models.values():
    kf = KFold(n_splits=6, shuffle=True, random_state=10)
    cv_results = cross_val_score(model, X_train, y_train, cv=kf)
    results.append(cv_results)

```

```

In [14]: plt.style.use('seaborn-v0_8')
fig, ax = plt.subplots(figsize=(16, 12))


boxplot = ax.boxplot(
    results,
    labels=models.keys(),
    patch_artist=True,
    boxprops=dict(edgecolor='black'),
    whiskerprops=dict(linewidth=1),
    capprops=dict(linewidth=1),
    medianprops=dict(color='black', linewidth=1),
    widths=0.2,
)

box_colors = ['#fceb03', '#03fcd9', '#90EE90', '#fc9d03', '#FFC0CB']
for box, color in zip(boxplot['boxes'], box_colors):
    box.set(facecolor=color)

ax.set_title('Wyniki walidacji krzyżowej', fontsize=24)
ax.tick_params(axis='both', labelsize=17)
ax.set_xlabel('Model', fontsize=20)
ax.set_ylabel('Precyzja', fontsize=20)
ax.set_ylim(0.70, 1)
ax.yaxis.grid(True)

```

```
plt.tight_layout()
plt.show()
```



Nauczenie **najskuteczniejszego modelu**

- Najlepszy okazał się **RandomForestClassifier**

```
In [15]: import sklearn.metrics as skm
random_forest = RandomForestClassifier(max_depth=None,
                                       min_samples_leaf=1,
                                       min_samples_split=10,
                                       n_estimators=50)

rand_forest.fit(X_train, y_train)
y_pred = rand_forest.predict(X_test)
```

Ponowne **sprawdzenie skuteczności modelu**

Tym razem dokładniej, przy pomocy:

- macierzy pomyłek** (confusion matrix)
- classification report**

```
In [16]: # Confusion Matrix
confusion_matrix = skm.confusion_matrix(y_test, y_pred)
```

```

plt.style.use('default')
cm_display = skm.ConfusionMatrixDisplay(
    confusion_matrix=confusion_matrix,
    display_labels=['safe', 'phishing'])


cm_display.plot(cmap='Blues')
plt.show()


# Wartości %
cm_display = skm.ConfusionMatrixDisplay(
    confusion_matrix=confusion_matrix/np.sum(confusion_matrix),
    display_labels=['safe', 'phishing'])

cm_display.plot(cmap='Blues')
plt.show()

# Classification report
target_names = ['safe', 'phishing']
class_report = skm.classification_report(y_test, y_pred, target_names=target_names)
print(class_report)

```





	precision	recall	f1-score	support
safe	0.87	0.90	0.89	3137
phishing	0.94	0.92	0.93	5178
accuracy			0.91	8315
macro avg	0.90	0.91	0.91	8315
weighted avg	0.91	0.91	0.91	8315

```
In [ ]: # Confusion Matrix
confusion_matrix = skm.confusion_matrix(y_test, y_pred)

plt.style.use('default')
fig, ax = plt.subplots(ncols=2, nrows=1)
ax[0] = skm.ConfusionMatrixDisplay(
    confusion_matrix=confusion_matrix,
    display_labels=['safe', 'phishing'])

ax[0].plot(cmap='Blues')

# Wartości %
ax[1] = skm.ConfusionMatrixDisplay(
    confusion_matrix=confusion_matrix/np.sum(confusion_matrix),
    display_labels=['safe', 'phishing'])

ax[1].plot(cmap='Blues')
plt.show()
```

```
# Classification report
target_names = ['safe', 'phishing']
class_report = skm.classification_report(y_test, y_pred, target_names=target_names)
print(class_report)
```

Tworzymy metodę do konwertowania maila na wartości liczbowe na wzór tych użytych przy nauczaniu modelu

```
In [23]: from unicode import unicode
import spacy
import re

nlp = spacy.load('en_core_web_md', disable=["parser", "ner"])
def import_suspicious_words():
    with open('suspicious_words.txt', 'r') as file:
        lines = file.readlines()

    suspicious_words = []
    for line in lines:
        word = line.strip()
        suspicious_words.append(word)

    suspicious_str = ' '.join(suspicious_words)
    doc = nlp(suspicious_str)
    suspicious_lemmas = [token.lemma_.lower() for token in doc]

    suspicious_set = set(suspicious_lemmas)
    return suspicious_set
def parse_mail_to_nums(mail: dict):

    sender = mail['sender']
    splitted = sender.split('@')

    sender_name = splitted[0]
    sender_domain = splitted[1]

    sender_num_count = sum(1 for char in sender_name if char.isnumeric())

    splitted_dom = sender_domain.split('.')[:-1]
    domain_noTLD = '.'.join(splitted_dom)
    print(domain_noTLD)
    sender_domain_num_count = sum(1 for char in domain_noTLD if char.isnumeric())
    print('sender_domain_num_count ', sender_domain_num_count)
    sender_domain_length = len(domain_noTLD)

    # subject and body text
    suspicious_set = import_suspicious_words()

    doc = nlp(mail['subject'])
    tokens = [token.lemma_.lower() for token in doc if token.text.isalnum()]
    unique_words = set(tokens)
    sus_words = unique_words.intersection(suspicious_set)
    suspicious_words_subject = len(sus_words)

    doc = nlp(mail['body'])
```

```

tokens = [token.lemma_.lower() for token in doc if token.text.isalnum()]
unique_words = set(tokens)
sus_words = unique_words.intersection(suspicious_set)
suspicious_words_body = len(sus_words)

# body urls
body = mail['body']
url_pattern = re.compile(r'https?://\S+|www\.\S+')
urls = re.findall(url_pattern, body)
# urls_count
urls_count = len(urls)
url = urls[np.random.randint(0, urls_count)]
# protocol
protocol = url[:5].lower()
protocol = 'https' if protocol == 'https' else 'http'
# contains_ip
IP_pattern = re.compile(r'\b(?:\d{1,3}\.){3}\d{1,3}\b|\b(?:[0-9a-fA-F]{1,4}:){7}')
IPs = IP_pattern.findall(url)
contains_ip = 1 if IPs else 0
# url_length
url_length = len(url)
# TLD_alpha
pattern = re.compile(r'https?://(?:[^\?]+)')
match = pattern.match(url)
if match:
    domain = match.group(1)
    if '/' in domain:
        domain = domain.split('/')[0]
else:
    domain = url

split_domain = domain.split('.')
n = len(split_domain)
delimiters = ['/', ':', ')', ']', '%', '_', '=', ',', '>', '"', '#', '!']
# Check if not weird ending
after_dot = split_domain[n-1]
if len(after_dot) > 2 and not after_dot.isalpha():
    after_dot = after_dot.split('/')[0]

    if len(after_dot) > 2 and not after_dot.isalpha():
        print(after_dot, 'INSIDE')
        for delimiter in delimiters:
            after_dot = " ".join(after_dot.split(delimiter))

    after_dot = after_dot.split()[0]

TLD = '.'+after_dot.lower()
TLD_alpha = TLD[1:].isalpha()
# subdomain_level
subdomain_level = domain.count('.')-1
# slash_count
slash_count = url.count('/')
# dots_count
dots_count = url.count('.')
# hyphens_count

```

```

hyphens_count = url.count('-')
# has_non_latin
ascii = unidecode(url)
has_non_latin = url != ascii

data = {
    'sender_num_count': sender_num_count,
    'sender_domain_num_count': sender_domain_num_count,
    'sender_domain_length': sender_domain_length,
    'suspicious_words_subject': suspicious_words_subject,
    'suspicious_words_body': suspicious_words_body,
    'urls_count': urls_count,
    'protocol': protocol,
    'contains_ip': contains_ip,
    'url_length': url_length,
    'TLD_alpha': TLD_alpha,
    'subdomain_level': subdomain_level,
    'slash_count': slash_count,
    'dots_count': dots_count,
    'hyphens_count': hyphens_count,
    'has_non_latin': has_non_latin,
}

for k, v in data.items():
    print('{k}: {v}'.format(k, v), end='\n')

# numeric values
urls_count_out = urls_count if urls_count <= 2 else 3
protocol_out = 1 if protocol=='https' else 0
contains_ip_out = contains_ip
url_length_out = url_length
TLD_alpha_out = 1 if TLD_alpha is True else 0
subdomain_level_out = subdomain_level if subdomain_level <= 2 else 3
slash_count_out = slash_count if slash_count <= 5 else 6
dots_count_out = dots_count if dots_count <= 4 else 5
hyphens_count_out = hyphens_count if hyphens_count <= 1 else 2
has_non_latin_out = has_non_latin

X_output = np.array(list([
    [sender_num_count],
    [sender_domain_num_count],
    [sender_domain_length],
    [suspicious_words_subject],
    [suspicious_words_body],
    [urls_count_out],
    [protocol_out],
    [contains_ip_out],
    [url_length_out],
    [TLD_alpha_out],
    [subdomain_level_out],
    [slash_count_out],
    [dots_count_out],
    [hyphens_count_out],
    [has_non_latin_out]
]))
return X_output.reshape(1, -1)

```

Testy na życiowych przykładach

```
In [50]: mail_example = {
'sender': 'Hanna.Wdowicka@ue.poznan.pl',
'subject': 'Reaktywacja SKN Estymator',
'body': '''
    Szanowni Państwo,

    Poniżej przekazuję wiadomość od dr Macieja Beęsewicza, prof. UEP. Osoby zai
    Pozdrawiam,
    Hanna Wdowicka

    Temat: Reaktywacja SKN Estymator

    Treść:

    Szanowni Państwo,

    dr Maciej Beręsewicz z Katedry Statystyki poszukuje studentów/tek zainteres
    Poniżej przedstawiono przykłady prac, które będą realizowane w ramach pracy

    1. Estimating the number of entities with vacancies using administrative an
    2. Inferring job vacancies from online job advertisements -- https://ec.eur
    3. Enhancing the Demand for Labour survey by including skills from online j
    4. Maximum entropy classification for record linkage – https://htmlpreview.
    5. Developing methods for determining the number of unauthorized foreigners

    Proszę się kontaktować przez mail maciej.beresewicz@ue.poznan.pl lub przez

    Z poważaniem

    dr Maciej Beręsewicz, prof. UEP

    Katedra Statystyki
    Uniwersytet Ekonomiczny w Poznaniu
    Department of Statistics
    Poznań University of Economics and Business
    al. Niepodległości 10 | 61-875 Poznań
    tel. + 48 61-854-36-80 | maciej.beresewicz@ue.poznan.pl
    www.ue.poznan.pl
    ...
    }
prediction = rand_forest.predict(
    scaler.transform(parse_mail_to_nums(mail_example))
)

print('-----\nPredykcja:\n-----')
print(colored('Phishing! >:(, 'red')) if prediction == 1 else print(colored('E-mai
```

```

ue.poznan
sender_domain_num_count 0
sender_num_count: 0
sender_domain_num_count: 0
sender_domain_length: 9
suspicious_words_subject: 0
suspicious_words_body: 2
urls_count: 6
protocol: https
contains_ip: 0
url_length: 52
TLD_alpha: True
subdomain_level: 2
slash_count: 6
dots_count: 3
hyphens_count: 1
has_non_latin: False
-----
Predykcja:
-----
E-mail bezpieczny 8)

```

```

In [32]: mail_example = {
    'sender': 'x3r0@ysadna321.com',
    'subject': 'Your xero invoice available now.',
    'body': '''
        Hi,
        Thanks for working with us. Your bill for $373.75 was due on 28 Aug 2016.
        If you've already paid it, please ignore this email and sorry for bothering
        To view your bill visit http://in.x312412.qwe12/5LQDhRwfvoQfeDtLDMqkk1JWSqC
        If you've got any questions, or want to arrange alternative payment don't h
        Thanks
        NJW Limited
        Download PDF
    '''
}
prediction = rand_forest.predict(
    scaler.transform(parse_mail_to_nums(mail_example))
)

print('-----\Predykcja:\n-----')
print(colored('Phishing! >:(, 'red')) if prediction == 1 else print(colored('E-mai

```

```
ysadna321
sender_domain_num_count 3
qwe12 INSIDE
sender_num_count: 2
sender_domain_num_count: 3
sender_domain_length: 9
suspicious_words_subject: 1
suspicious_words_body: 2
urls_count: 1
protocol: http
contains_ip: 0
url_length: 65
TLD_alpha: False
subdomain_level: 1
slash_count: 3
dots_count: 3
hyphens_count: 0
has_non_latin: False
-----\Predykcja:
-----
Phishing! >:(
```

✦ ✦ *Rezultat projektu* ✦ ✦

Przy pomocy **RandomForestClassifier** udało się wytrenować model z wynikami:

- **accuracy:** 0.91
- **precision:**
 - safe: 0.87
 - phishing: 0.94
- **recall:**
 - safe: 0.90
 - phishing: 0.92
- **f1:**
 - safe: 0.89
 - phishing: 0.93

Problemy projektu

- mało **uniwersalny**
- **model** za bardzo **opiera się** na **odnośnikach** znajdujących się w e-mailu
 - bardzo dużo maili **phishingowych** zawiera **buttony**, nie **klasyczne odnośniki**
- **dostęp** do jakiegokolwiek sensownego **API** jest **płatny**
 - brak sprawdzania **domen w blacklistach**
 - brak sprawdzania adresów **e-mail w blacklistach**
 - słabe/brak informacji o **szyfrowaniu SSL**
 - brak sprawdzania **wieku** domeny

- Dane **zbierane** w latach **1998-2022**
 - prawdopodobnie wiele maili pochodzi z wczesnych lat 2000
 - (bardzo dużo http, nawet wśród bezpiecznych domen)
 - Człowiek jakkolwiek obeznany w internecie bez problemu poradziłby sobie z klasyfikowaniem ataków **phishingowych** z datasetu.
-